Detecting Insults in Social Commentary

IMPERMIUM

Predict whether a comment posted during a public discussion is considered insulting to one of the participants.

$10,000 · 50 teams · 5 years ago

Competition Data

Edit

- impermium_verificati...
- impermium_verificati...
- sample_submission_nu...
- test.csv
- test_with_solutions....
- train.csv

train.csv 830.53 KB

Download

Data Description

Data

The data consists of a label column followed by two attribute fields.

This is a single-class classification problem. The label is either 0 meaning a **neutral** comment, or 1 meaning an **insulting** comment (neutral can be considered as not belonging to the insult class. Your predictions must be a real number in the range [0,1] where 1 indicates 100% confident prediction that comment is an insult.

The first attribute is the time at which the comment was made. It is sometimes blank, meaning an accurate timestamp is not possible. It is in the form "YYYYMMDDhhmmss" and then the Z character. It is on a 24 hour clock and corresponds to the localtime at which the comment was originally made.

The second attribute is the unicode-escaped text of the content, surrounded by double-quotes. The content is mostly english language comments, with some occasional formatting.

Guidelines

- We are looking for comments that are intended to be insulting to a person who is a part of the larger blog/forum conversation.
- We are NOT looking for insults directed to non-participants (such as celebrities, public figures etc.).
- Insults could contain profanity, racial slurs, or other offensive language. But often times, they do not.
- Comments which contain profanity or racial slurs, but are not necessarily insulting to another person are considered not insulting.
- The insulting nature of the comment should be obvious, and not subtle.

There may be a small amount of noise in the labels as they have not been meticulously cleaned. However, contenstants can be confident the error in the training and testing data is < 1%.
Contestants should also be warned that this problem tends to strongly overfit. The provided data is generally representative of the full test set, but not exhaustive by any measure. Impermium will be conducting final evaluations based on an unpublished set of data drawn from a wide sample.